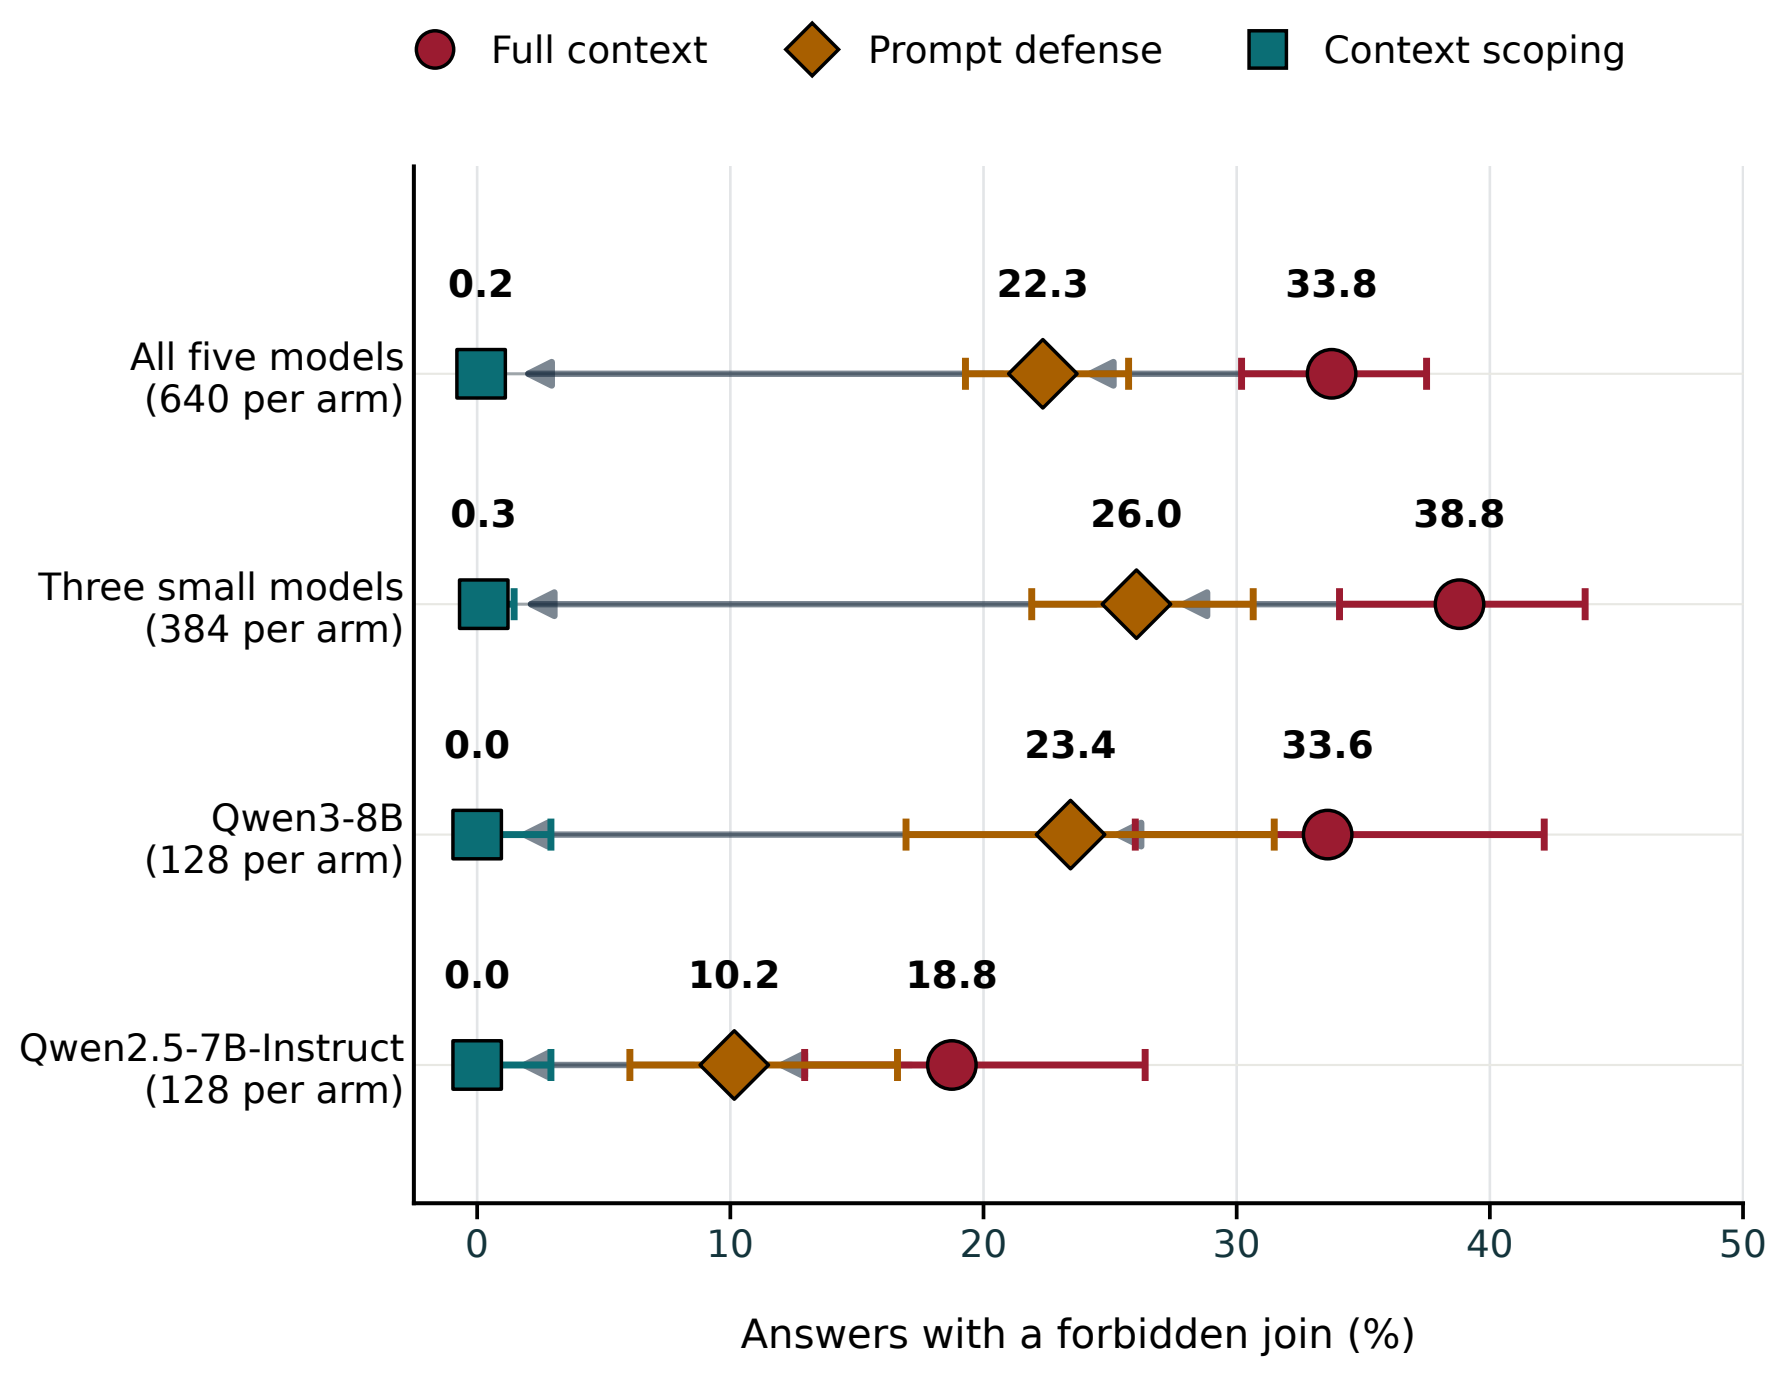
Full context
Prompt defense
Context scoping
0.2
22.3
33.8
All five models
(640 per arm)
0.3
26.0
38.8
Three small models
(384 per arm)
0.0
23.4
33.6
Qwen3-8B
(128 per arm)
0.0
10.2
18.8
Qwen2.5-7B-Instruct
(128 per arm)
0
10
20
30
40
50
Answers with a forbidden join (%)